LLM Evaluation Report

SageMaker JumpStart models

Generated: 2026-09-28 20:08:22
Test Models: sagemaker/huggingface-od-nvidia-locateanything-3b
Judge Model: none (geometric IoU against exact ground truth; prompts: prompt_templates_grounding.json)
AWS Region: us-east-1
DeepEval: Disabled
Endpoints: huggingface-od-nvidia-locateanything-3b → jumpstart-dft-hf-od-nvidia-locatean-20260821-170355

Executive Summary

Model	Tests	Pass Rate	Avg Score	Avg Latency
sagemaker/huggingface-od-nvidia-locat...	3	66.7%	0.66	1284ms

Performance by Category

Category	Tests	Pass Rate	Avg Score	Status
ATTRIBUTE_BINDING	3	66.7%	0.66	WARNING

Model: sagemaker/huggingface-od-nvidia-locateanything-3b

Category: ATTRIBUTE_BINDING

Test 1: Colour Discrimination, Same Shape - PASS | Score: 0.98

Input Prompt:

```
Locate all the instances that matches the following description: blue circle.
```

AI Response:

```
Model ref: blue circle. the Distinct boxes returned: 1 (repeated 2x in raw output) Expected boxes: 1
expected px=[380, 140, 520, 280] -> predicted px=[378, 139, 522, 282] IoU=0.952 (match)
```

Method	Score	Status
Localization	0.95	PASS
Detection	1.00	PASS
Format	1.00	PASS

Observations: localization: 0.95 | detection: 1.00 | format: 1.00 | mean IoU 0.952 over 1 target(s) | probe: Two circles differing only in colour. Must bind colour to shape.

Recommendations: Accurate localization and count. No action needed.

Tokens: 0->0 | Latency: 712ms | Time: 2026-09-28 20:08:19

Test 2: Shape Discrimination, Same Colour - FAIL | Score: 0.15

Input Prompt:

```
Locate all the instances that matches the following description: red square.
```

AI Response:

```
Model ref: (none) Distinct boxes returned: 0 (repeated 0x in raw output) Expected boxes: 1
expected px=[380, 140, 520, 280] -> MISSED
```

Method	Score	Status
Localization	0.00	FAIL
Detection	0.00	FAIL
Format	1.00	PASS

Observations: localization: 0.00 | detection: 0.00 | format: 1.00 | mean IoU 0.000 over 1 target(s) | 1 missed | probe: Two red objects differing only in shape.

Recommendations: Returned no boxes for a target that is present. Try rephrasing the description (this model is phrasing-sensitive) before concluding the object is undetectable.

Tokens: 0->0 | Latency: 332ms | Time: 2026-09-28 20:08:19

Test 3: Shape Discrimination, Rephrased Query - PASS | Score: 0.85

Input Prompt:

```
Locate all the instances that matches the following description: red rectangle.
```

AI Response:

```
Model ref: red rectangle. the Distinct boxes returned: 2 (repeated 21x in raw output) Expected
boxes: 1 expected px=[380, 140, 520, 280] -> predicted px=[380, 139, 522, 281] IoU=0.972
(match) extra prediction px=[401, 0, 640, 480] -> false positive
```

Method	Score	Status
Localization	0.97	PASS
Detection	0.67	WARNING
Format	1.00	PASS

Observations: localization: 0.97 | detection: 0.67 | format: 1.00 | mean IoU 0.972 over 1 target(s) | 1 false positive(s) | output repeated 21x (token waste) | probe: Identical scene to attr_shape_red with one word changed. Probes phrasing sensitivity: the endpoint returned 0 boxes for 'red square' but located the square for 'red rectangle' on 2026-09-28.

Recommendations: Count is wrong (missed targets or false positives). Review instance separation for this scene type.

Tokens: 0->0 | Latency: 2807ms | Time: 2026-09-28 20:08:22